



Let's Play a Game!



- **Knowns:**
 - 4 possible keys (1-4)
 - Reward-based
 - Starting position
- **Unknowns:**
 - Goal of the game
 - What any of the keys does
 - Consequences of objects in the environment

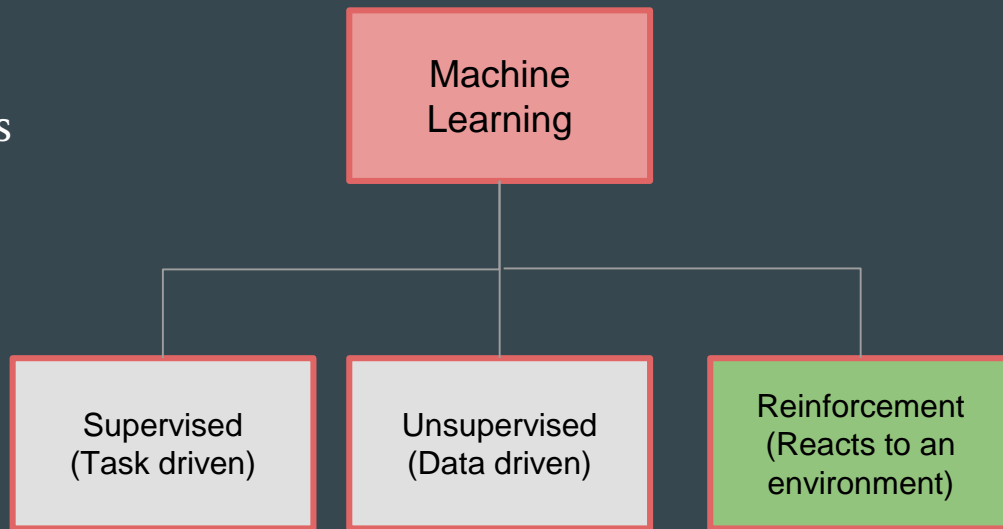
Deep Q-Learning



Brittany Haffner, Atharva Jakkanwar, Apar Singhal, Bryan Fuh

Reinforcement Learning

- Neither supervised or unsupervised
- Does not require training data
- Interacts with the environment
- Updates policy based on rewards



Markov Decision Process

- Framework used to make decision on stochastic environment
- More useful than simple planning
 - MDP: solution will give you optimal action even if something goes wrong
 - Simple planning: Follow the plan after you find best strategy (doesn't allow mistakes/errors)
- Markovian property: the effects of an action taken in a state depend only on that state and not on prior history

Markov Decision Process

- Discrete time stochastic control process:
 - States: s
 - Actions: $A(s), a$
 - Rewards: $R(s), R(s,a)$
 - E.g. -1 for every move
 - E.g. +100/-100 for end states
 - Transition model (rules/physics of the game): $T(s, a, s') \sim Pr(s' / s, a)$
 - For this example, 0.8 chance of going in the desired direction
 - Policy (solution): $\pi(s) \rightarrow a; \pi^*(s) \rightarrow a$

			Goal
			Die
Start			

Markov Decision Process

- Infinite horizon
- Finite horizon
 - Finite number of time steps
 - Changes policy
- Utility of sequences: $U(s_0, s_1, s_2, \dots) = \sum_t \gamma^t R(s_t)$, $0 \leq \gamma \leq 1$
 - γ : discount factor
 - Allows program to go infinite distance in finite amount of time

			Goal
			Die
			Start

Bellman Equations

- The key to solving MDP in RL
- $\pi^* = \operatorname{argmax}_{\pi} \mathbb{E}[\sum_t \gamma^t R(s_t) | \pi] \rightarrow \pi^*(s) = \operatorname{argmax}_a \sum_{s'} T(s, a, s') U(s')$
- Utility score
 - The expected maximum cumulative reward by following the policy given by state s
 - $U(s) = \alpha [R(s, a) + \gamma \max_a \sum_{s'} T(s, a, s') U(s')]$, $0 < \alpha < 1$ (learning rate)
 - $\pi^*(s) = \operatorname{argmax}_a U(s, a)$
 - $U(s) = R(s) + \gamma \pi^*(s)$
- Fundamental recursive equation that defines the true value in a particular state
- Determines the best action to take

Bellman Equations

- Quality score (Q-Value)

- The expected maximum cumulative reward from taking action a in state s and following the policy
- $Q(s,a) = \alpha\{R(s,a) + \gamma \max_{a'} \sum_{s'} T(s, a, s') Q(s', a')\}$
- $Q(s', a') \leftarrow Q(s,a) + \alpha\{R(s,a) + \gamma \max_{a'} Q(s', a') - Q(s,a)\}$
- New action value = Old value + learning rate * (new information - old information)

- Finding policies:

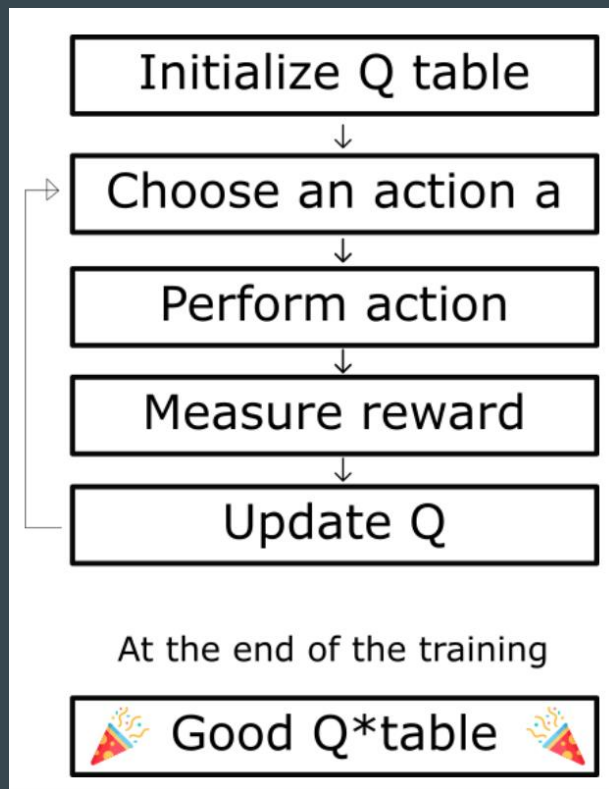
- Initialize: π_0 as initial guess
- Evaluate: given π_t calculate $Q_t = Q_t^\pi$
- Improve: $\pi_{t+1} = \operatorname{argmax}_a \sum T(s, a, s') Q_t(s', a')$
- Policy equation: $Q_t(s) = \alpha\{R(s) + \gamma \sum_{s'} T(s, \pi_t(t), s') Q_t(s')\}$

Q-Learning

- Q-value-based reinforcement learning
- Take action based on maximum expected future reward for current state: $\pi(s) = \operatorname{argmax}_a Q(s, a)$
- Generate a Q-table

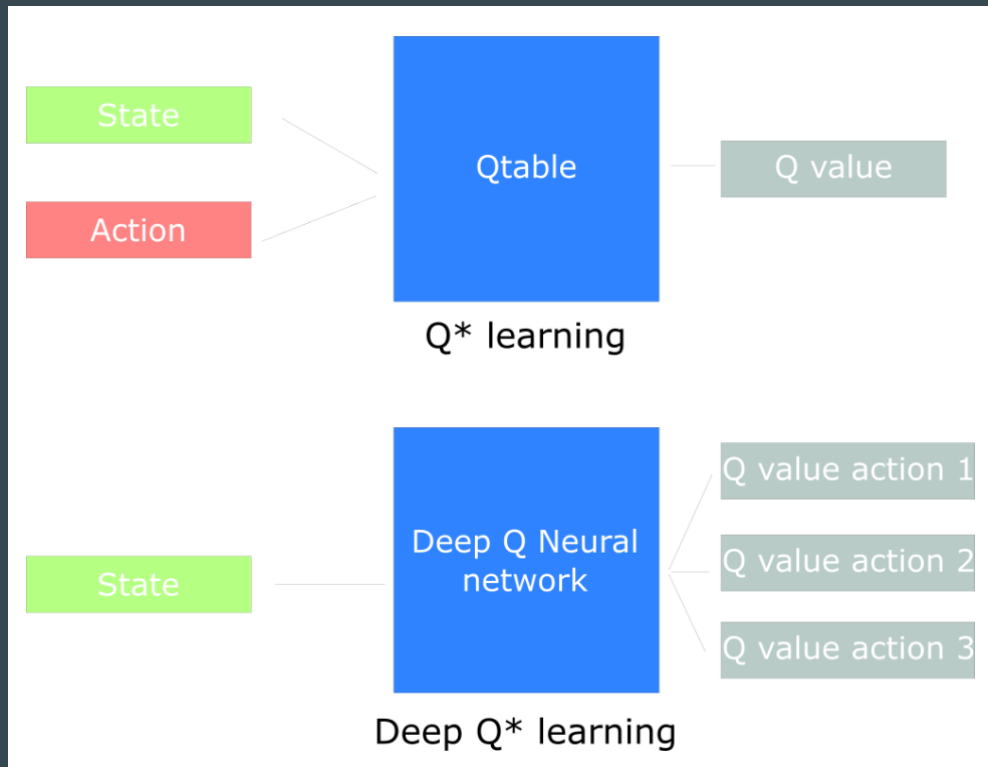
Example Q-table

	a_1	a_2	$\dots a_n$
s_1	0	0.1	0
s_2	0	0	0
\dots	0	0	0
s_m			



Deep Q-Learning vs. Q-Learning

- Replace Q-table with CNN
- Better for large environments with many different states
- Used for learning control policies from high-dimensional inputs
- Choose the action with the highest Q-value

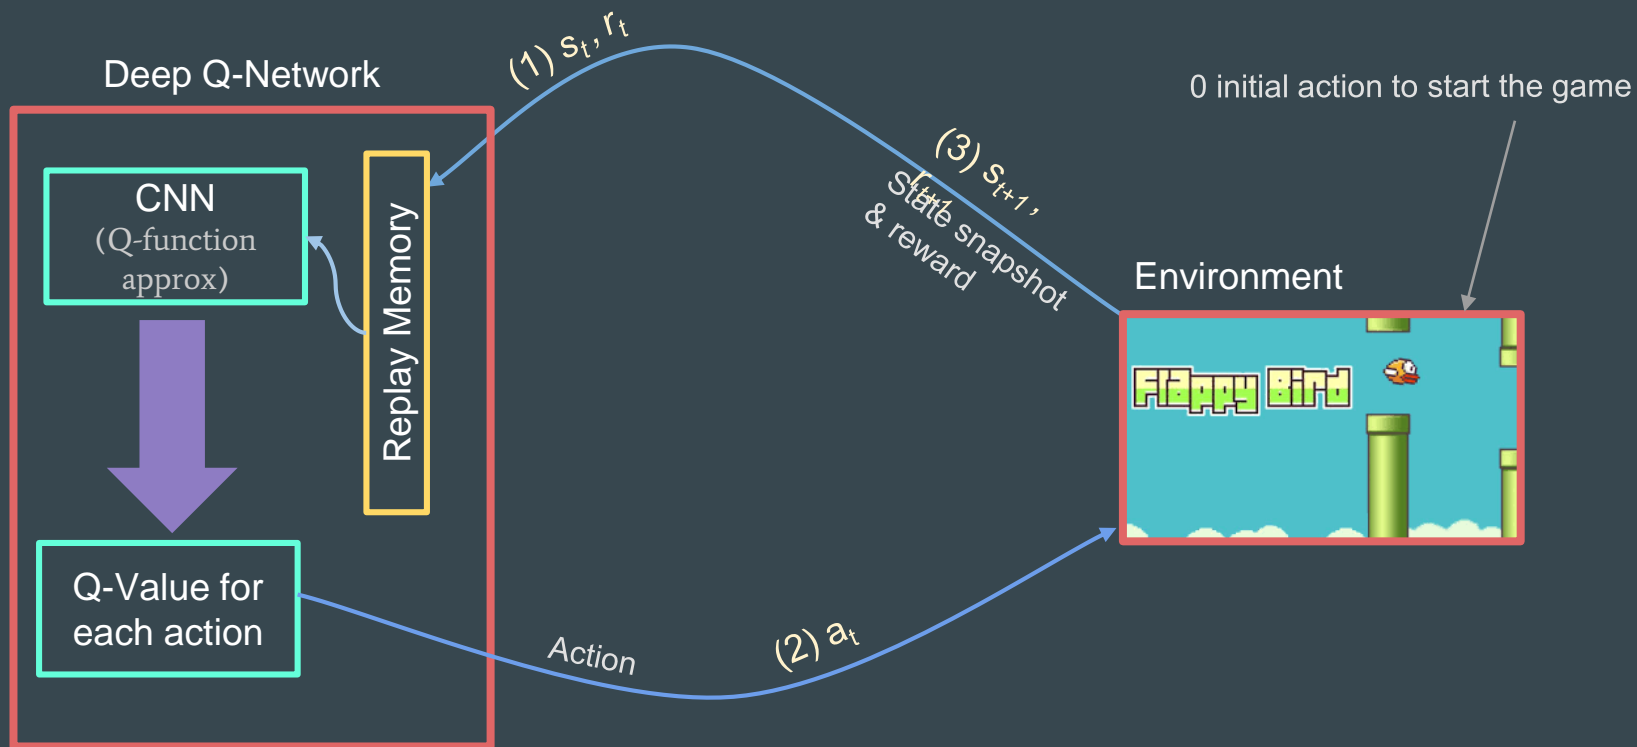


Problem Statement

- Environment: Flappy Bird (pygame)
- Goal: Continually keeping the bird alive while it's moving forward
- State (s): Position of the bird in the environment (determined by the physics)
- Actions (a): Tap the screen or don't tap the screen
- Rewards (r):
 - +0.1 every timestep it lives
 - +1 every pipe it crosses
 - -1 every time it dies (terminal state)
 - Reward clipping for better generalization

(MDP - perform an action a_t at state s_t to receive reward r_t and next state s_{t+1})


Architecture



Architecture (CNN)

- Preprocessing images for faster learning/training
 - RGB -> Grayscale
 - Resize to a smaller size
 - Stack 4 frames

$$L = \mathbb{E}[\underbrace{(r + \gamma \max_{a'} Q(s', a') - Q(s, a))^2}_{\text{target}}]$$




C(8x8x4x32)A -> MP(2x2) -> C(4x4x32x64)A -> C(3x3x64x64)A -> FC(1600)A -> FC(512) -> FC(2)

Epsilon Greedy Policy (Exploration vs Exploitation)

- To learn the optimal policy π^* , the agent must see all the possible states
 - So it should be exposed to as many states as possible
 - But the agent has access to only states that are generated by its own action
- Greedy approach (max Q) is optimal within the space already explored
- Epsilon provides new info to the agent; the probability of taking a random action
- Will make decisions based in new data obtained through randomization
- Mostly epsilon in annealed in order to reduce the exploration

Experience Replay

- In reinforcement learning, we receive sequential samples from interactions with the environment, which are highly correlated and can lead to overfitting
- Experience Replay to the rescue!
 - Avoid forgetting previous experiences
 - Reduce correlations between experiences
 - Distribution averaged over many previous states, smoothing out the learning
- Used to train Deep-Q Network (DQN)



Deep Q-Learning with Experience Replay

Initialize environment E

Initialize replay memory M with finite capacity N

Initialize the DQN weights w

for $episode$ in $max_episode$:

for $steps$ in max_steps :

s = Environment state

Choose action a from state s using probability ϵ greedily


Take action a , get reward R and next state s'

Store experience $\langle s_t, a_t, r_t, s_{t+1} \rangle$ in M


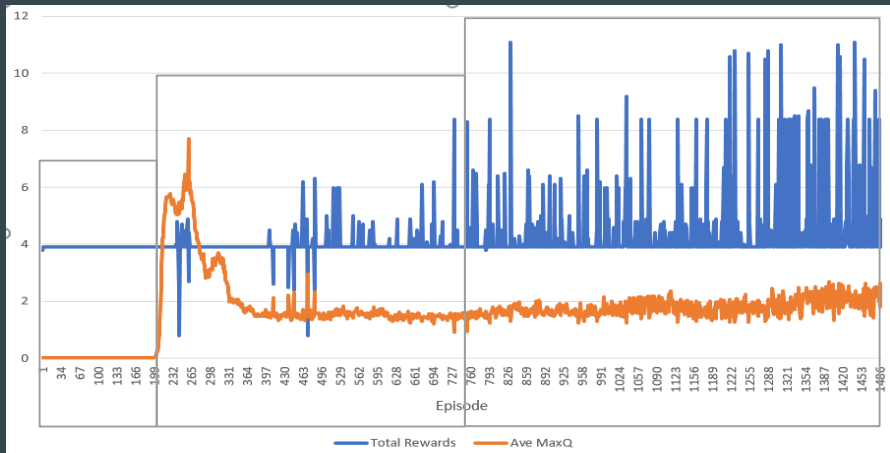

Get random minibatch of experience $\langle s_p, a_p, r_p, s_{p'} \rangle$ from M

Set optimal Q-value $Q(s_p)^* = r_p + \gamma \max Q(s_{p'})$, where γ is discount value

Update Q by taking gradient descent step on $(Q(s_p)^* - Q(s_p))^2$



Flappy Bird Demonstration



Thank you!

References:

- [Stanford Lecture: Reinforcement Learning](#)
- [Human-level control through deep reinforcement learning](#)
- [Playing Atari with Deep Reinforcement Learning](#)
- [Udacity: Reinforcement Learning GaTech](#)
- [Pygame implementation of Flappy Bird](#)